

Multiple Disease Prediction Model

Tamanna Sharma

Lovely Professional University

12110221

CONTENTS

Sl. No.	Topic	Page No.
1	Acknowledgement	03
2	Project Objective	04-05
3	Project Scope	06
4	Data Description	07-08
5	Data Pre-Processing	09-17
6	Model Building	18-25
7	Test Dataset	26-29
8	Code	30-79
9	Future Scope of Improvements	80
10	Certificates	81

ACKNOWLEDGEMENT

I take this opportunity to express my profound gratitude and deep regards to my faculty, **Prof. Arnab Chakraborty** for his exemplary guidance, monitoring and constant encouragement throughout the course of this project. The blessing, help and guidance given by him time to time shall carry me a long way in the journey of life on which I am about to embark.

I extend my gratitude to the authors and researchers whose work I have referenced throughout this project. Their groundbreaking research, published papers, and open-source contributions have laid the foundation upon which my project is built. I am indebted to their dedication and commitment to advancing the field of machine learning.

Project Objective

In this project we have made two models, 'Diabetes Prediction Model' and 'Heart Disease Prediction Model' and we have made a web app with a menu bar having options to predict any of them.

Let us first discuss about the Diabetes Dataset which we have taken from Kaggle. In this dataset, the target attribute is outcome which is binary. So, in this project we need to do binary classification based on the attributes present in our dataset and predict whether a person is suffering from diabetes or not.

Similarly, Heart Disease Dataset is also taken from Kaggle. In this dataset, the target attribute is target which is also binary. Therefore, we need to do the binary classification for heart disease prediction model as well.

Our objective in this project is to make a web app using stream-lit library that will be able to predict the diabetes and heart disease in an individual. For that we first need to create ML models for each one of them. We need to study each dataset and we might need to pre-process the given dataset if we need to. Then, we would train 4 models viz. 'KNN classifier model', 'Naive Bayes classifier model', 'Support Vector Machine Model' and 'Logistic Regression model'. After training the aforementioned models, we will need to find out the score and classification report. Our next step would be to use the trained models to predict the outcomes using the given test dataset and compare the outcome of each model. We would then choose the best model based on the accuracy score and classification report. After this using pickle5 library we will be dumping the model and when required we will load the model and using stream-lit we will be creating option menu bar and a proper and complete web app to take inputs from users and provide the predicted output in a better way.

Our methodology for solving the problems in the given project is described below:

- Load the required dataset.
- Study the dataset.
- Describe the dataset.
- Visualise the dataset.
- Find out if the dataset needs to be pre-processed. It will be determined on the basis of whether the dataset has null values or outliers or any such discrepancy that might affect the output of the models to be trained.
- If the dataset is required to be pre-processed, take the necessary steps to pre-process the data.

- Find out the principal attributes for training.
- Split the given dataset for training and testing purpose.
- Fit the previously split train data in the aforementioned 4 models.
- Calculate the accuracy of the 4 models and find out the classification reports.
- Plot the necessary graphs.
- Use each trained model to predict the outcomes of the given test dataset.
- Choose the best model among the 4 trained models bases on the accuracy and classification reports and build the web app using stream-lit.

PROJECT SCOPE

The broad scope of 'Multiple Disease Prediction Model' project is given below:

- The given datasets have the attributes based on which it can be predicted that whether a person is suffering from a particular disease or not.
- It is a useful project as the Classifier models can be used to quickly determine whether an individual is healthy or not.
- These models can be used in hospitals by the doctors and they can modify them according to their needs. This will reduce the manual labour and time spent on determining whether a person has a particular disease or not.
- It can also provide personalized predictions to patients.
- The dataset given to us is a shortened form of the original dataset from Kaggle. So, the results might have some mismatch with the real-world applications. But that can be
- avoided if the models are trained accordingly.

DATA DESCRIPTION

Source of the data: Kaggle. The given dataset is a shortened version of the original dataset in Kaggle.

Diabetes Dataset: This dataset has 768 rows and 8 columns.

- Pregnancies: Number of times pregnant
- Glucose: Plasma glucose concentration a 2 hours in an oral glucose tolerance test
- BloodPressure: Diastolic blood pressure (mm Hg)
- SkinThickness: Triceps skin fold thickness (mm)
- Insulin: 2-Hour serum insulin (mu U/ml)
- BMI: Body mass index (weight in kg/(height in m)^2)
- DiabetesPedigreeFunction: Diabetes pedigree function
- Age: Age (years)
- Outcome: Class variable (0 or 1)

Class Distribution: (class value 1 is interpreted as "tested positive for diabetes")

Heart Disease Dataset: This dataset has 303 rows and 13 columns.

1. age
2. sex
3. cp: chest pain type (4 values)
4. trestbps: resting blood pressure
5. chol: serum cholestrol in mg/dl
6. fbs: fasting blood sugar > 120 mg/dl
7. restecg: resting electrocardiographic results (values 0,1,2)
8. thalach: maximum heart rate achieved
9. exang: exercise induced angina
10. oldpeak: ST depression induced by exercise relative to rest
11. slope: the slope of the peak exercise ST segment
12. ca: number of major vessels (0-3) colored by flourosopy
13. thal: 0 = normal; 1 = fixed defect; 2 = reversable defect
14. target: Class variable (0 or 1). It is integer valued 0 = no disease and 1 = disease.

The following table shows 5 number statistics of the diabetes dataset:

	Pregnancies	Glucose	Blood Pressure	Skin Thickness	Insulin	BMI	Diabetes Pedigree Function	Age	Outcome
mean	3.845052	120.894531	69.105469	20.666667	79.799479	31.992578	0.471876	32.890625	0.348958
std	3.369578	31.972618	19.355807	15.561237	115.24400	7.884160	0.331329	11.500287	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.25000	36.600000	0.626250	40.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

Similarly, we can also know about the statistics of heart disease dataset which would be given in code part.

Now we will pre-process the data. The methodology followed is given below:

- Checking for null values.
 - o If null values are present, we will fill them or drop the row containing the null value based on the dataset.
- Checking for outliers.
 - o If outliers are present, they will either be removed or replaced by following a suitable method depending on the dataset.


DATA PRE-PROCESSING

We searched for null values in our dataset and formed the following table:

Diabetes:

Pregnancies	0
Glucose	0
BloodPressure	0
SkinThickness	42
Insulin	0
BMI	0
DiabetesPedigreeFunction	0
Age	45
Outcome	0


To visualise the null values, we made a heatmap plot using seaborn library function heatmap. The heatmap plot is given below:



The heatmap shows that the dataset has null values.

To remove the null values we have filled the missing values with the median in age attribute and SkinThickness attribute.

After removing the null values, the following heatmap was obtained:




Similarly, the null values are removed from heart disease dataset:

Count of null values:

Age	27
Sex	0
Cp	0
Trestbps	0
Chol	0
Fbs	0
Restecg	0
Thalach	23
Exang	0
Oldpeak	0
Slope	0
Ca	0
Thal	0
Target	0


To visualise the null values, we made a heatmap plot using seaborn library function heatmap. The heatmap plot is given below:



The heatmap shows that the dataset has null values.

To remove the null values we have filled the missing values with the median in age attribute and thalach attribute.

After removing the null values, the following heatmap was obtained:



Now we have successfully handled Null values and we didn't drop the rows with null values as we have a small dataset.

So, we are moving on to find if there are any outliers in our data and find the correlations of different attributes to our target i.e. 'outcome' in diabetes and 'target' in heart disease dataset.

Pregnancies	0.221898
Glucose	0.466581
BloodPressure	0.065068
SkinThickness	0.074752
Insulin	0.130548
BMI	0.292695
DiabetesPedigreeFunction	0.173844
Age	0.238356
Outcome	1.000000

Diabetes dataset correlation with Outcome

Age	-0.217072
Sex	-0.280937
Cp	0.433798
Trestbps	-0.144931
Chol	-0.085239
Fbs	-0.028046
Restecg	0.137230
Thalach	0.403064
Exang	-0.436757
Oldpeak	-0.430696
Slope	0.345877
Ca	-0.391724
Thal	-0.344029
Target	1.000000



Heart Disease dataset correlation with target



Outliers are extreme values that deviate from other observations on data, they may indicate a variability in a measurement, experimental errors or a novelty. In other words, an outlier is an observation that diverges from an overall pattern on a sample.


Most common causes of outliers on a data set:

- Data entry errors (human errors)
- Measurement errors (instrument errors)
- Experimental errors (data extraction or experiment planning/executing errors)
- Intentional (dummy outliers made to test detection methods)
- Data processing errors (data manipulation or data set unintended mutations)
- Sampling errors (extracting or mixing data from wrong or various sources)
- Natural (not an error, novelties in data)

We plot distribution graph to visualise the outliers in diabetes dataset. The plots are given below:



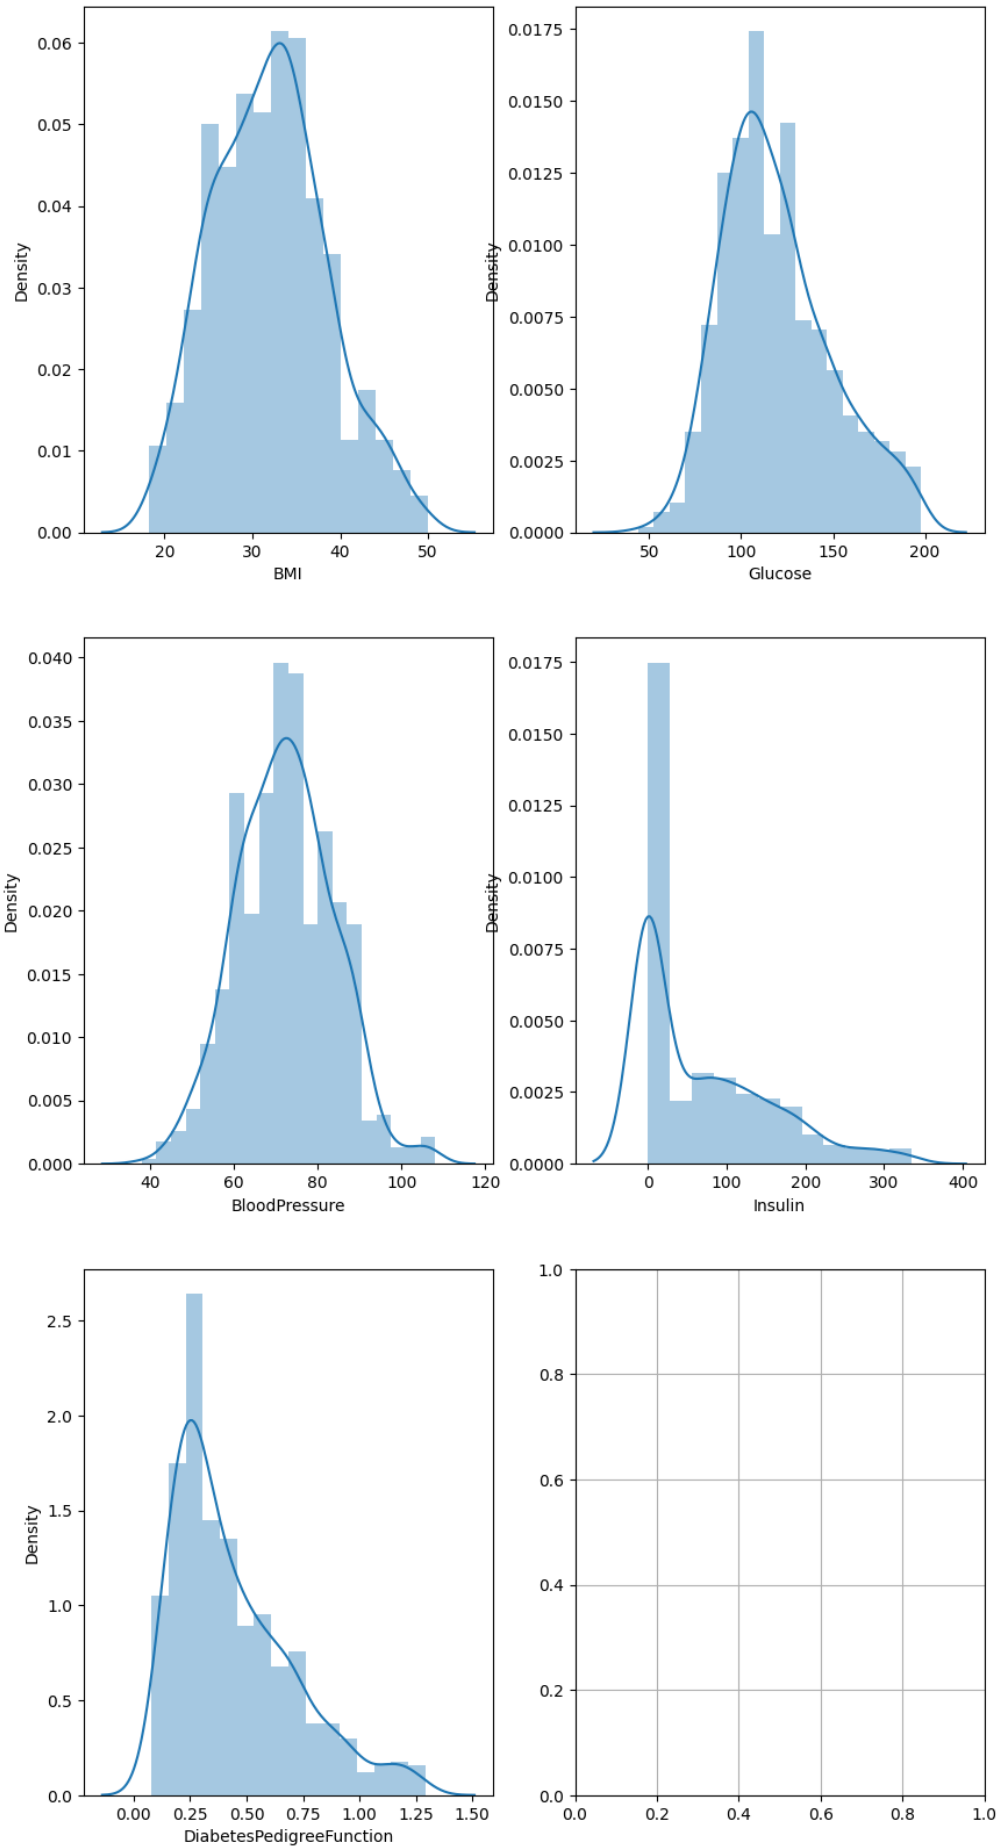




We have removed outliers using Z scores. We have created a function which will return a series where True represents a row that contains an outlier (absolute z-score greater than 3) and False does not.

We have also created a function that will iterate over all the columns and will provide us a dataframe denoting whether there is an outlier or not and lastly we will be filtering out the outliers.

After removing outliers:



MODEL BUILDING

Splitting data for training and testing purpose


We split the given train dataset into two parts for training and testing purpose. The split ratio we used is 0.80 which indicates we used 80% data for training purpose and 20% data for testing purpose. We will be using the same split ratio for all the models trained. The same models are used for both the prediction systems. We will be looking at the diabetes system here.

SVM classifier

Support Vector Machine or SVM is one of the most popular Supervised Learning algorithms, which is used for Classification as well as Regression problems. However, primarily, it is used for Classification problems in Machine Learning.

The goal of the SVM algorithm is to create the best line or decision boundary that can segregate n-dimensional space into classes so that we can easily put the new data point in the correct category in the future. This best decision boundary is called a hyperplane.

SVM chooses the extreme points/vectors that help in creating the hyperplane. These extreme cases are called as support vectors, and hence algorithm is termed as Support Vector Machine. Consider the below diagram in which there are two different categories that are classified using a decision boundary or hyperplane:



Confusion Matrix:

Now we created a confusion matrix to view the actual and predicted test results. Given below is the confusion matrix:

Predicted	0	1
Actual		
0	82	8
1	23	20

Classification Report:

	precision	recall	f1-score	support
0	0.78	0.91	0.84	90
1	0.71	0.47	0.56	43
accuracy			0.77	133
macro avg	0.75	0.69	0.70	133
weighted avg	0.76	0.77	0.75	133


Logistic Regression model:

- Logistic regression is one of the most popular Machine Learning algorithms, which comes under the Supervised Learning technique. It is used for predicting the categorical dependent variable using a given set of independent variables.
- Logistic regression predicts the output of a categorical dependent variable. Therefore the outcome must be a categorical or discrete value. It can be either Yes or No, 0 or 1, true or False, etc. but instead of giving the exact value as 0 and 1, **it gives the probabilistic values which lie between 0 and 1.**
- Logistic Regression is much similar to the Linear Regression except that how they are used. Linear Regression is used for solving Regression problems, whereas **Logistic regression is used for solving the classification problems.**

- The sigmoid function is a mathematical function used to map the predicted values to probabilities.
- It maps any real value into another value within a range of 0 and 1.
- The value of the logistic regression must be between 0 and 1, which cannot go beyond this limit, so it forms a curve like the "S" form. The S-form curve is called the Sigmoid function or the logistic function.
- In logistic regression, we use the concept of the threshold value, which defines the probability of either 0 or 1. Such as values above the threshold value tends to 1, and a value below the threshold values tends to 0.

Assumptions for Logistic Regression:

- The dependent variable must be categorical in nature.
- The independent variable should not have multi-collinearity.



Confusion matrix:

Predicted	0	1
Actual		
0	81	9
1	22	21


Classification Report:

	precision	recall	f1-score	support
0	0.79	0.90	0.84	90
1	0.70	0.49	0.58	43
accuracy			0.77	133
macro avg	0.74	0.69	0.71	133
weighted avg	0.76	0.77	0.75	133

KNN Classifier:

k-NN can be used for both classification and regression predictive problems. However, it is more widely used in classification problems in the industry. In pattern recognition, the k-nearest neighbors algorithm (k-NN) is a non-parametric method used for classification and regression. In both cases, the input consists of the k closest training examples in the feature space. The output depends on whether k-NN is used for classification or regression. In k-NN classification, the output is a class membership. An object is classified by a plurality vote of its neighbors, with the object being assigned to the class most common among its k nearest neighbors (k is a positive integer, typically small). For e.g. if k = 1, then the object is simply assigned to the class of that single nearest neighbor. k-NN is a type of instance-based learning, or lazy learning, where the function is only approximated locally and all computation is deferred until classification. The k-NN algorithm is among the simplest of all

machine learning algorithms. The neighbors are taken from a set of objects for which the class (for k-NN classification) or the object property value (for k-NN regression) is known. This can be thought of as the training set for the algorithm, though no explicit training step is required.



The K-NN working can be explained on the basis of the below algorithm:

- **Step-1:** Select the number K of the neighbors
- **Step-2:** Calculate the Euclidean distance of **K number of neighbors**
- **Step-3:** Take the K nearest neighbors as per the calculated Euclidean distance.
- **Step-4:** Among these k neighbors, count the number of the data points in each category.
- **Step-5:** Assign the new data points to that category for which the number of the neighbor is maximum.
- **Step-6:** Our model is ready.

Confusion Matrix:

Predicted	0	1
Actual		

0	83	7
1	25	18

Classification report:

	precision	recall	f1-score	support
0	0.77	0.92	0.84	90
1	0.72	0.42	0.53	43
accuracy			0.76	133
macro avg	0.74	0.67	0.68	133
weighted avg	0.75	0.76	0.74	133

Naïve Bayes Model:

A Naive Bayes classifier is a probabilistic machine learning model that's used for classification task. The crux of the classifier is based on the Bayes theorem. The fundamental Naive Bayes assumption is that each feature makes an:

- independent
- equal

contribution to the outcome. The assumptions made by Naive Bayes are not generally correct in real-world situations. In-fact, the independence assumption is never correct but often works well in practice.

Bayes Theorem:

Bayes' Theorem finds the probability of an event occurring given the probability of another event that has already occurred. Bayes' theorem is stated mathematically as the following equation:

$$P(A|B) = P(A)P(B|A)/P(B)$$

where A and B are events and P(B) is the probability of occurrence of event B.

Basically, we are trying to find probability of event A, given the event B is true. Event B is also termed as evidence. P(A) is the priori of A (the prior probability, i.e. Probability of event before evidence is seen). The evidence is an attribute value of an unknown instance (here, it

is event B). $P(A|B)$ is a posteriori probability of B, i.e. probability of event after evidence is seen. The class-data relation from the Bayes Theorem can be obtained as follows:

$$P(\text{Class} | \text{Data}) = P(\text{Class})P(\text{Data} | \text{Class})/P(\text{Data})$$

Where,

- $P(\text{Class} | \text{Data})$ = Posterior
- $P(\text{Class})$ = Prior
- $P(\text{Data} | \text{Class})$ = Likelihood
- $P(\text{Data})$ = Marginal Probability

In other words, it can be written as:

$$\text{Posterior} = \text{Prior} * \text{Likelihood} / \text{Marginal Probability}$$

In application, we do not need to calculate the Marginal Probability for classification. We only need to calculate the numerator of the posterior for classification.

Types of Naive Bayes Classifier:

Multinomial Naive Bayes:

This is mostly used for document classification problem, i.e. whether a document belongs to the category of sports, politics, technology etc. The features/predictors used by the classifier are the frequency of the words present in the document.

Bernoulli Naive Bayes:

This is similar to the multinomial Naive Bayes but the predictors are Boolean variables. The parameters that we use to predict the class variable take up only values yes or no, for example if a word occurs in the text or not.

Gaussian Naive Bayes:

When the predictors take up a continuous value and are not discrete, we assume that these values are sampled from a gaussian distribution.

Confusion Matrix:

Predicted	0	1
Actual		
0	72	18
1	21	22

Classification report:

	precision	recall	f1-score	support
0	0.77	0.80	0.79	90
1	0.55	0.51	0.53	43
accuracy			0.71	133
macro avg	0.66	0.66	0.66	133
weighted avg	0.70	0.71	0.70	133

Similarly, all the models are trained in heart disease prediction system.

Comparison of the Models trained:


We trained 4 models using the 4 algorithms viz.

1. k-Nearest Neighbour
2. Gaussian Naive Bayes
3. SVM
4. Logistic Regression

The 4 models had different accuracy.

For Diabetes Prediction:


The comparison of the accuracies of the models are given below(**without outliers**):



Thus, from the above comparison we can see that the KNN model has the highest accuracy. So, our selected model is KNN Model.

Model	Accuracy %
k-Nearest Neighbour	78.195
Gaussian Naive Bayes	69.924
SVM	76.691
Logistic Regression	75.939

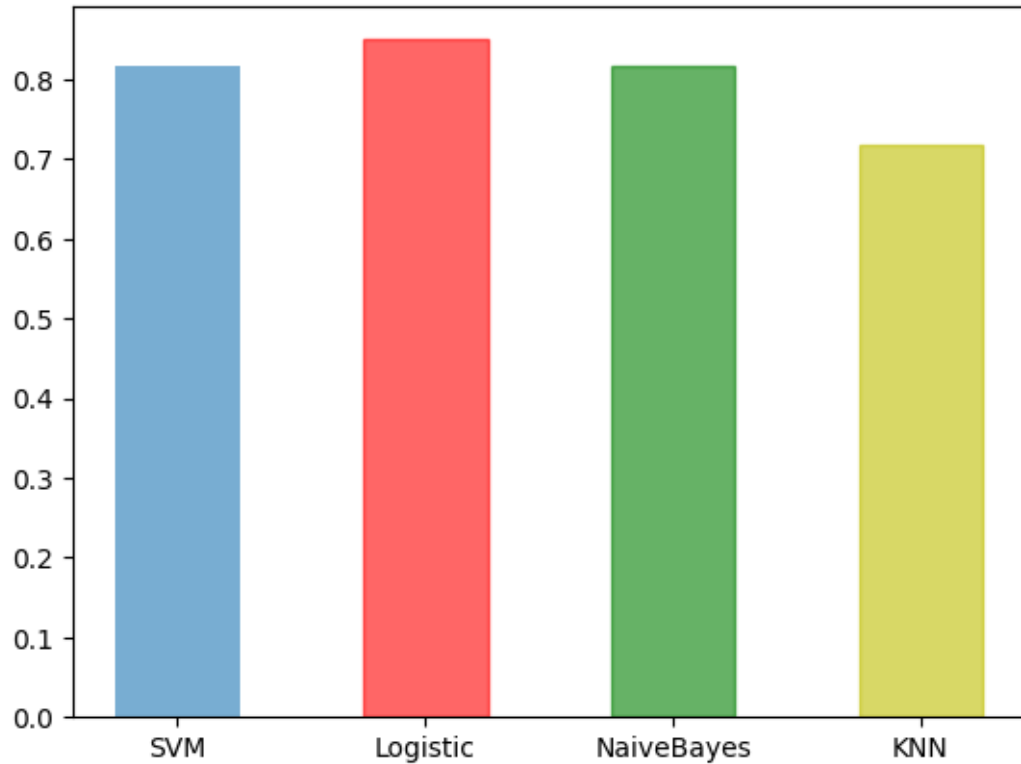
With outliers:



Model	Accuracy %
k-Nearest Neighbour	72.077
Gaussian Naive Bayes	77.272
SVM	77.272
Logistic Regression	75.974

For Heart Disease Prediction:


The comparison of the accuracies of the models are given below(**without outliers**):



Thus, from the above comparison we can see that the Logistic Regression model has the highest accuracy. So, our selected model is Logistic Regression Model.

Model	Accuracy %
k-Nearest Neighbour	71.667
Gaussian Naive Bayes	81.667
SVM	81.667
Logistic Regression	85

The comparison of the accuracies of the models (**with outliers**):



Model	Accuracy %
k-Nearest Neighbour	67.213
Gaussian Naive Bayes	80.327
SVM	80.327
Logistic Regression	78.688

CODES

MULTIPLE DISEASE PREDICTION SYSTEM:

DIABETES PREDICTION SYSTEM:

Importing the dependencies

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
```

Data Collection And Analysis PIMA Diabetes Dataset

```
#Loading diabetes dataset to pandas dataframe
df=pd.read_csv('/content/diabetes.csv')
diabetes_df=df.copy()
```

diabetes_df

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35.0	0	33.6	
1	1	85	66	29.0	0	26.6	
2	8	183	64	NaN	0	23.3	
3	1	89	66	23.0	94	28.1	
4	0	137	40	35.0	168	43.1	
..	
763	10	101	76	48.0	180	32.9	
764	2	122	70	27.0	0	36.8	
765	5	121	72	23.0	112	26.2	
766	1	126	60	0.0	0	30.1	
767	1	93	70	31.0	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50.0	1
1	0.351	31.0	0
2	0.672	NaN	1
3	0.167	21.0	0
4	2.288	33.0	1
..
763	0.171	63.0	0
764	0.340	27.0	0
765	0.245	30.0	0
766	0.349	47.0	1
767	0.315	23.0	0

[768 rows x 9 columns]

```
diabetes_df.isnull().sum()
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness    25
Insulin           0
BMI              0
DiabetesPedigreeFunction  0
Age              71
Outcome          0
dtype: int64
```

```
diabetes_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Pregnancies	768 non-null	int64
1	Glucose	768 non-null	int64
2	BloodPressure	768 non-null	int64
3	SkinThickness	743 non-null	float64
4	Insulin	768 non-null	int64
5	BMI	768 non-null	float64
6	DiabetesPedigreeFunction	768 non-null	float64
7	Age	697 non-null	float64
8	Outcome	768 non-null	int64

```
dtypes: float64(4), int64(5)
```


```
memory usage: 54.1 KB
```

```
#heatmap
```

```
import seaborn as sns
```

```
sns.heatmap(diabetes_df.isnull(),yticklabels=False,cbar=True,cmap="viridis")
```

```
<Axes: >
```



#distribution of data in age

```
fig,ax=plt.subplots(figsize=(8,8))
sns.distplot(diabetes_df.Age)
```

<ipython-input-512-a3786ae00d7b>:3: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.


Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.Age)
```



```
<Axes: xlabel='Age', ylabel='Density'>
```



```
sns.distplot(diabetes_df.SkinThickness)
```

```
<ipython-input-513-2fcc0f9ea930>:1: UserWarning:
```


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.SkinThickness)
```

```
<Axes: xlabel='SkinThickness', ylabel='Density'>
```



```
sns.distplot(diabetes_df.Pregnancies)
```

```
<ipython-input-514-6d99d7eb254e>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.Pregnancies)
```

```
<Axes: xlabel='Pregnancies', ylabel='Density'>
```



```
sns.distplot(diabetes_df.Glucose)
```

```
<ipython-input-515-6aef5e1f755c>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.Glucose)
```

```
<Axes: xlabel='Glucose', ylabel='Density'>
```



```
sns.distplot(diabetes_df.BloodPressure)
```

```
<ipython-input-516-9385a01ab079>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.BloodPressure)
```

```
<Axes: xlabel='BloodPressure', ylabel='Density'>
```



```
sns.distplot(diabetes_df.Insulin)
```

```
<ipython-input-517-80c25523b1df>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn
v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.Insulin)
```

```
<Axes: xlabel='Insulin', ylabel='Density'>
```



```
sns.distplot(diabetes_df.BMI)
```

<ipython-input-518-d512f55bcf04>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.BMI)
```

<Axes: xlabel='BMI', ylabel='Density'>



```
sns.distplot(diabetes_df.DiabetesPedigreeFunction)
```

<ipython-input-519-d282ce82f0d2>:1: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.DiabetesPedigreeFunction)
```

<Axes: xlabel='DiabetesPedigreeFunction', ylabel='Density'>



#replace the missing values with median values

```
diabetes_df['Age'].fillna(diabetes_df['Age'].median(),inplace=True)
diabetes_df['SkinThickness'].fillna(diabetes_df['SkinThickness'].median(),
inplace=True)
```

```
def find_outliers(col):
    from scipy import stats
    z=np.abs(stats.zscore(col))
    idx_outliers=np.where(z>3,True,False)
    return pd.Series(idx_outliers,index=col.index)
idx=find_outliers(diabetes_df.BMI)
idx1=find_outliers(diabetes_df.Glucose)
idx2=find_outliers(diabetes_df.BloodPressure)
idx3=find_outliers(diabetes_df.Insulin)
idx4=find_outliers(diabetes_df.DiabetesPedigreeFunction)
idx
```

```
0      False
1      False
2      False
3      False
4      False
...
763    False
764    False
765    False
766    False
```



```
767     False
Length: 768, dtype: bool
```

```
display(diabetes_df.loc[idx==True].describe().round(3))
display(diabetes_df.loc[idx1==True].describe().round(3))
display(diabetes_df.loc[idx2==True].describe().round(3))
display(diabetes_df.loc[idx3==True].describe().round(3))
display(diabetes_df.loc[idx4==True].describe().round(3))
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
count	14.000	14.000	14.000	14.000	14.000	14.000
mean	3.286	112.786	43.214	13.571	33.786	13.129
std	3.384	27.096	46.161	21.103	71.595	26.167
min	0.000	74.000	0.000	0.000	0.000	0.000
25%	0.000	96.000	0.000	0.000	0.000	0.000
50%	2.500	114.500	32.000	0.000	0.000	0.000
75%	5.750	124.500	81.000	23.000	10.500	0.000
max	10.000	180.000	110.000	63.000	240.000	67.100

	DiabetesPedigreeFunction	Age	Outcome
count	14.000	14.000	14.000
mean	0.599	29.143	0.286
std	0.673	14.250	0.469
min	0.102	21.000	0.000
25%	0.238	22.000	0.000
50%	0.304	24.500	0.000
75%	0.623	26.000	0.750
max	2.420	69.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
count	5.00	5.0	5.000	5.000	5.000	5.000
mean	2.80	0.0	67.600	30.200	4.600	32.880
std	2.49	0.0	12.033	8.643	10.286	7.034
min	1.00	0.0	48.000	20.000	0.000	24.700
25%	1.00	0.0	68.000	23.000	0.000	27.700
50%	1.00	0.0	68.000	32.000	0.000	32.000
75%	5.00	0.0	74.000	35.000	0.000	39.000
max	6.00	0.0	80.000	41.000	23.000	41.000

	DiabetesPedigreeFunction	Age	Outcome
count	5.000	5.000	5.000
mean	0.380	27.800	0.400
std	0.216	6.221	0.548
min	0.140	22.000	0.000
25%	0.299	22.000	0.000
50%	0.346	29.000	0.000
75%	0.389	29.000	1.000
max	0.727	37.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
count	35.000	35.000	35.0	35.000	35.0	35.000

mean	3.629	117.800	0.0	1.514	0.0	25.706
std	3.647	27.489	0.0	6.298	0.0	14.875
min	0.000	73.000	0.0	0.000	0.0	0.000
25%	0.000	97.500	0.0	0.000	0.0	21.650
50%	2.000	117.000	0.0	0.000	0.0	28.900
75%	6.000	133.500	0.0	0.000	0.0	34.550
max	13.000	183.000	0.0	30.000	0.0	52.300

	DiabetesPedigreeFunction	Age	Outcome
count	35.000	35.000	35.000
mean	0.388	29.257	0.457
std	0.254	6.294	0.505
min	0.102	21.000	0.000
25%	0.198	25.000	0.000
50%	0.282	29.000	0.000
75%	0.573	31.000	1.000
max	0.933	44.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
count	18.000	18.000	18.000	18.000	18.000	18.000
mean	3.111	165.833	73.444	35.333	548.833	36.961
std	2.888	20.057	11.638	9.107	107.813	5.988
min	0.000	124.000	52.000	23.000	440.000	28.700
25%	1.000	155.000	63.500	26.250	480.000	31.375
50%	2.000	168.500	72.000	35.500	502.500	37.550
75%	4.750	180.000	83.500	43.500	570.500	40.300
max	8.000	197.000	90.000	49.000	846.000	52.300

	DiabetesPedigreeFunction	Age	Outcome
count	18.000	18.000	18.000
mean	0.661	34.722	0.667
std	0.625	13.663	0.485
min	0.128	21.000	0.000
25%	0.244	23.500	0.000
50%	0.540	29.500	1.000
75%	0.687	45.250	1.000
max	2.329	60.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
count	11.000	11.000	11.000	11.000	11.000	11.000
mean	1.909	140.909	67.273	30.182	170.000	33.055
std	2.468	37.920	13.154	16.940	240.219	15.139
min	0.000	82.000	40.000	0.000	0.000	0.000
25%	0.000	118.000	61.000	21.000	0.000	25.950
50%	2.000	137.000	70.000	24.000	89.000	36.700
75%	2.500	176.500	77.000	39.000	221.500	41.250
max	8.000	197.000	82.000	63.000	744.000	59.400

	DiabetesPedigreeFunction	Age	Outcome
count	11.000	11.000	11.000
mean	1.914	28.273	0.364
std	0.324	6.198	0.505

min	1.476	21.000	0.000
25%	1.699	25.000	0.000
50%	1.781	25.000	0.000
75%	2.212	30.000	1.000
max	2.420	44.000	1.000

```
display(diabetes_df.loc[idx==False].describe().round(3))
diabetes_df=diabetes_df.loc[idx==False]
display(diabetes_df.loc[idx1==False].describe().round(3))
diabetes_df=diabetes_df.loc[idx1==False]
display(diabetes_df.loc[idx2==False].describe().round(3))
diabetes_df=diabetes_df.loc[idx2==False]
display(diabetes_df.loc[idx3==False].describe().round(3))
diabetes_df=diabetes_df.loc[idx3==False]
display(diabetes_df.loc[idx4==False].describe().round(3))
diabetes_df=diabetes_df.loc[idx4==False]
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
count	754.000	754.000	754.000	754.000	754.000
754.000					
mean	3.855	121.045	69.586	21.077	80.654
32.343					
std	3.371	32.052	18.224	15.548	115.756
6.690					
min	0.000	0.000	0.000	0.000	0.000
18.200					
25%	1.000	99.000	64.000	0.000	0.000
27.500					
50%	3.000	117.000	72.000	23.000	36.500
32.250					
75%	6.000	141.000	80.000	32.000	128.750
36.600					
max	17.000	199.000	122.000	99.000	846.000
55.000					

	DiabetesPedigreeFunction	Age	Outcome
count	754.000	754.000	754.000
mean	0.470	32.798	0.350
std	0.322	11.241	0.477
min	0.078	21.000	0.000
25%	0.244	24.000	0.000
50%	0.376	29.000	0.000
75%	0.626	39.000	1.000
max	2.329	81.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
count	749.000	749.000	749.000	749.000	749.000
749.000					
mean	3.862	121.853	69.599	21.016	81.162
32.339					
std	3.376	30.588	18.263	15.569	115.972
6.692					

min	0.000	44.000	0.000	0.000	0.000
18.200					
25%	1.000	99.000	64.000	0.000	0.000
27.500					
50%	3.000	117.000	72.000	23.000	38.000
32.300					
75%	6.000	141.000	80.000	32.000	130.000
36.600					
max	17.000	199.000	122.000	99.000	846.000
55.000					

	DiabetesPedigreeFunction	Age	Outcome
count	749.000	749.000	749.000
mean	0.470	32.832	0.350
std	0.323	11.262	0.477
min	0.078	21.000	0.000
25%	0.244	24.000	0.000
50%	0.376	29.000	0.000
75%	0.626	39.000	1.000
max	2.329	81.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
count	721.000	721.000	721.000	721.000	721.000
721.000					
mean	3.878	121.791	72.302	21.759	84.313
32.347					
std	3.363	30.737	12.281	15.336	117.075
6.642					
min	0.000	44.000	24.000	0.000	0.000
18.200					
25%	1.000	99.000	64.000	8.000	0.000
27.500					
50%	3.000	117.000	72.000	23.000	48.000
32.300					
75%	6.000	142.000	80.000	33.000	130.000
36.600					
max	17.000	199.000	122.000	99.000	846.000
55.000					

	DiabetesPedigreeFunction	Age	Outcome
count	721.000	721.000	721.000
mean	0.472	32.922	0.343
std	0.325	11.404	0.475
min	0.078	21.000	0.000
25%	0.245	24.000	0.000
50%	0.378	29.000	0.000
75%	0.626	39.000	1.000
max	2.329	81.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
count	703.000	703.000	703.000	703.000	703.000
703.000					

mean	3.898	120.663	72.273	21.411	72.420
32.229					
std	3.374	30.136	12.304	15.309	90.011
6.619					
min	0.000	44.000	24.000	0.000	0.000
18.200					
25%	1.000	99.000	64.000	0.000	0.000
27.400					
50%	3.000	116.000	72.000	23.000	43.000
32.000					
75%	6.000	139.000	80.000	32.000	125.500
36.400					
max	17.000	199.000	122.000	99.000	415.000
55.000					

	DiabetesPedigreeFunction	Age	Outcome
count	703.000	703.000	703.000
mean	0.467	32.876	0.334
std	0.313	11.348	0.472
min	0.078	21.000	0.000
25%	0.245	24.000	0.000
50%	0.375	29.000	0.000
75%	0.619	39.000	1.000
max	2.288	81.000	1.000

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
BMI \					
count	696.000	696.000	696.000	696.000	696.000
696.000					
mean	3.917	120.609	72.359	21.398	72.346
32.224					
std	3.376	30.126	12.260	15.341	89.881
6.602					
min	0.000	44.000	24.000	0.000	0.000
18.200					
25%	1.000	99.000	64.000	0.000	0.000
27.400					
50%	3.000	115.000	72.000	23.000	43.500
32.050					
75%	6.000	139.000	80.000	32.000	125.250
36.325					
max	17.000	199.000	122.000	99.000	415.000
55.000					

	DiabetesPedigreeFunction	Age	Outcome
count	696.000	696.000	696.000
mean	0.454	32.907	0.335
std	0.284	11.383	0.472
min	0.078	21.000	0.000
25%	0.245	24.000	0.000
50%	0.370	29.000	0.000
75%	0.605	39.000	1.000
max	1.461	81.000	1.000

```

df_outliers=pd.DataFrame()
for col in df.describe().columns:
    df_outliers[col]=find_outliers(diabetes_df[col])
df_outliers.head()

```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	False	False	False	False	False	False	
1	False	False	False	False	False	False	
2	False	False	False	False	False	False	
3	False	False	False	False	False	False	
5	False	False	False	False	False	False	

	DiabetesPedigreeFunction	Age	Outcome
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
5	False	False	False

```

test_outs=df_outliers.apply(lambda x:np.any(x),axis=1)

print(len(test_outs),df_outliers.shape)
test_outs

696 (696, 9)

0      False
1      False
2      False
3      False
5      False
...
763     False
764     False
765     False
766     False
767     False
Length: 696, dtype: bool

np.sum(test_outs)

33

df_clean=diabetes_df.loc[test_outs==False]
df_clean.describe()
diabetes_df=df_clean

f,axes=plt.subplots(3,2,figsize=(10,20))
sns.distplot(diabetes_df.BMI,ax=axes[0][0])
plt.grid()
sns.distplot(diabetes_df.Glucose,ax=axes[0][1])
plt.grid()
sns.distplot(diabetes_df.BloodPressure,ax=axes[1][0])
plt.grid()

```

```
sns.distplot(diabetes_df.Insulin,ax=axes[1][1])
plt.grid()
sns.distplot(diabetes_df.DiabetesPedigreeFunction,ax=axes[2][0])
plt.grid()
```

<ipython-input-529-52febba1a34b>:2: UserWarning:

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.BMI,ax=axes[0][0])
<ipython-input-529-52febba1a34b>:4: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.Glucose,ax=axes[0][1])
<ipython-input-529-52febba1a34b>:6: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.BloodPressure,ax=axes[1][0])
<ipython-input-529-52febba1a34b>:8: UserWarning:
```

`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function

with
similar flexibility) or ``histplot`` (an axes-level function for
histograms).

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>







```
sns.distplot(diabetes_df.Insulin,ax=axes[1][1])  
<ipython-input-529-52febb1a34b>:10: UserWarning:
```

``distplot`` is a deprecated function and will be removed in seaborn
v0.14.0.

Please adapt your code to use either ``displot`` (a figure-level function
with
similar flexibility) or ``histplot`` (an axes-level function for
histograms).

For a guide to updating your code to use the new functions, please see
<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(diabetes_df.DiabetesPedigreeFunction,ax=axes[2][0])
```





```
diabetes_df.isnull().sum()
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```

```
sns.heatmap(diabetes_df.isnull(),yticklabels=False,cbar=True,cmap="viridis")
```

```
<Axes: >
```



```
diabetes_df.shape
```

```
(663, 9)
```

diabetes_df.corr()

	Pregnancies	Glucose	BloodPressure	
SkinThickness \				
Pregnancies	1.000000	0.154614	0.215797	-
0.092597				
Glucose	0.154614	1.000000	0.237370	
0.012857				
BloodPressure	0.215797	0.237370	1.000000	
0.042274				
SkinThickness	-0.092597	0.012857	0.042274	
1.000000				
Insulin	-0.094486	0.255382	-0.034030	
0.462058				
BMI	0.040603	0.196233	0.307080	
0.378410				
DiabetesPedigreeFunction	0.000952	0.066116	0.033880	
0.150030				
Age	0.528851	0.241782	0.333496	-
0.142780				
Outcome	0.241065	0.490638	0.190532	
0.045772				
	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.094486	0.040603	0.000952	
Glucose	0.255382	0.196233	0.066116	
BloodPressure	-0.034030	0.307080	0.033880	
SkinThickness	0.462058	0.378410	0.150030	
Insulin	1.000000	0.185041	0.195903	
BMI	0.185041	1.000000	0.131511	
DiabetesPedigreeFunction	0.195903	0.131511	1.000000	
Age	-0.065586	0.061163	0.041320	
Outcome	0.104582	0.281973	0.201347	
	Age	Outcome		
Pregnancies	0.528851	0.241065		
Glucose	0.241782	0.490638		
BloodPressure	0.333496	0.190532		
SkinThickness	-0.142780	0.045772		
Insulin	-0.065586	0.104582		
BMI	0.061163	0.281973		
DiabetesPedigreeFunction	0.041320	0.201347		
Age	1.000000	0.206625		
Outcome	0.206625	1.000000		


diabetes_df.corrwith(diabetes_df.Outcome)

Pregnancies	0.241065
Glucose	0.490638
BloodPressure	0.190532
SkinThickness	0.045772
Insulin	0.104582
BMI	0.281973
DiabetesPedigreeFunction	0.201347

```
Age                                0.206625
Outcome                           1.000000
dtype: float64
```

```
plt.figure(figsize = (10, 5))
sns.heatmap(diabetes_df.corr(), annot = True, annot_kws = {"size":15})
```

<Axes: >



#statistical measures of the data

```
diabetes_df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
\					
count	663.000000	663.000000	663.000000	663.000000	663.000000
mean	3.852187	119.583710	72.286576	21.135747	68.761689
std	3.300992	29.656877	11.497901	14.992535	84.273332
min	0.000000	44.000000	38.000000	0.000000	0.000000
25%	1.000000	99.000000	64.000000	0.000000	0.000000
50%	3.000000	114.000000	72.000000	23.000000	40.000000
75%	6.000000	137.500000	80.000000	32.000000	122.000000
max	14.000000	198.000000	108.000000	60.000000	335.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	663.000000	663.000000	663.000000	663.000000
mean	32.071493	0.441163	32.523379	0.324284
std	6.452334	0.266151	10.947098	0.468460
min	18.200000	0.078000	21.000000	0.000000
25%	27.350000	0.244000	24.000000	0.000000
50%	32.000000	0.365000	29.000000	0.000000
75%	36.100000	0.592000	38.000000	1.000000
max	50.000000	1.292000	67.000000	1.000000

```
diabetes_df.Outcome.value_counts() #diabetes_df['Outcome'].value_counts()
```

```
0    448
```

```
1    215
```

```
Name: Outcome, dtype: int64
```

0---->Non Diabetic 1----->Diabetic

```
diabetes_df.groupby('Outcome').mean()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
Outcome					
0	3.301339	109.511161	70.770089	20.660714	62.660714
1	5.000000	140.572093	75.446512	22.125581	81.474419

	BMI	DiabetesPedigreeFunction	Age
Outcome			
0	30.812054	0.404067	30.957589
1	34.695814	0.518460	35.786047

#Separating data and Labels

```
X=diabetes_df.drop('Outcome',axis=1)
```

```
Y=diabetes_df['Outcome']
```

```
X.head(5)
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35.0	0	33.6	
1	1	85	66	29.0	0	26.6	
2	8	183	64	23.0	0	23.3	
3	1	89	66	23.0	94	28.1	
5	5	116	74	0.0	0	25.6	

	DiabetesPedigreeFunction	Age
0	0.627	50.0
1	0.351	31.0
2	0.672	29.0
3	0.167	21.0
5	0.201	30.0

```
Y
```

```
0    1
1    0
2    1
3    0
5    0
..
763  0
764  0
765  0
766  1
767  0
```

```
Name: Outcome, Length: 663, dtype: int64
```

Data Standardization

```
scaler=StandardScaler()

scaler.fit(X)

StandardScaler()

standardized_data=scaler.transform(X)

standardized_data

array([[ 0.6511481 ,  0.95889205, -0.02494303, ...,  0.23707093,
         0.69876731,  1.59766692],
       [-0.86469175, -1.16700828, -0.54717137, ..., -0.84862689,
        -0.33902209, -0.13926328],
       [ 1.25748404,  2.13994779, -0.72124749, ..., -1.36045586,
         0.86797211, -0.32209803],
       ...,
       [ 0.34798013,  0.04779191, -0.02494303, ..., -0.91066676,
        -0.73759339, -0.23068066],
       [-0.86469175,  0.21651416, -1.06939972, ..., -0.30577798,
        -0.34654231,  1.32341478],
       [-0.86469175, -0.89705268, -0.19901915, ..., -0.25924808,
        -0.47438593, -0.87060231]])

X=standardized_data
Y=diabetes_df['Outcome']

print(X)
print(Y)

[[ 0.6511481  0.95889205 -0.02494303 ... 0.23707093 0.69876731
  1.59766692]
 [-0.86469175 -1.16700828 -0.54717137 ... -0.84862689 -0.33902209
 -0.13926328]
 [ 1.25748404  2.13994779 -0.72124749 ... -1.36045586  0.86797211
 -0.32209803]
 ...
 [ 0.34798013  0.04779191 -0.02494303 ... -0.91066676 -0.73759339
 -0.23068066]
 [-0.86469175  0.21651416 -1.06939972 ... -0.30577798 -0.34654231
  1.32341478]
 [-0.86469175 -0.89705268 -0.19901915 ... -0.25924808 -0.47438593
 -0.87060231]]
0      1
1      0
2      1
3      0
5      0
..
763    0
764    0
765    0
766    1
```

```
767      0
Name: Outcome, Length: 663, dtype: int64
```

Train Test Split

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,stratify=
Y,random_state=2)
```

```
print(X.shape,X_train.shape,X_test.shape)
```

```
(663, 8) (530, 8) (133, 8)
```

Support Vector Machine

```
#from sklearn.linear_model import LogisticRegression
classifier=svm.SVC(kernel="linear")
#classifier=LogisticRegression()
```

```
#training svmc
classifier.fit(X_train,Y_train)
```

```
SVC(kernel='linear')
```

Model Evaluation

```
#accuracy score on the training data
X_train_prediction=classifier.predict(X_train)
training_data_accuracy=accuracy_score(X_train_prediction,Y_train)
```

Printing accuracy score

```
print(training_data_accuracy)
```

```
0.7981132075471699
```

```
#accuracy score on the test data
X_test_prediction=classifier.predict(X_test)
svm_data_accuracy=accuracy_score(X_test_prediction,Y_test)
print(svm_data_accuracy)
print(classification_report(Y_test,X_test_prediction))
print(confusion_matrix(Y_test,X_test_prediction))
```

```
#
#          predicted
# actual   class-0   class-1
# class-0   TN       FP
# class-1   FN       TP
```

```
0.7669172932330827
```

	precision	recall	f1-score	support
0	0.78	0.91	0.84	90
1	0.71	0.47	0.56	43
accuracy			0.77	133
macro avg	0.75	0.69	0.70	133
weighted avg	0.76	0.77	0.75	133

```
[[82  8]
 [23 20]]
```

Logistic Regression

```
from sklearn.linear_model import LogisticRegression
clf=LogisticRegression()

clf.fit(X_train,Y_train)

LogisticRegression()

log_train_prediction=clf.predict(X_train)
print(accuracy_score(Y_train,log_train_prediction))

0.7981132075471699

log_test_prediction=clf.predict(X_test)
log_data_accuracy=accuracy_score(Y_test,log_test_prediction)
print(log_data_accuracy)

0.7593984962406015

print(classification_report(Y_test,log_test_prediction))
```

	precision	recall	f1-score	support
0	0.78	0.89	0.83	90
1	0.68	0.49	0.57	43
accuracy			0.76	133
macro avg	0.73	0.69	0.70	133
weighted avg	0.75	0.76	0.75	133

```
print(confusion_matrix(Y_test,log_test_prediction))

[[80 10]
 [22 21]]
```

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
rclf=RandomForestClassifier()

rclf.fit(X_train,Y_train)

RandomForestClassifier()

train_predictions=rclf.predict(X_train)
print(accuracy_score(Y_train,train_predictions))

1.0

test_predictions=rclf.predict(X_test)
random_data_accuracy=accuracy_score(Y_test,test_predictions)
print(random_data_accuracy)
```


0.7669172932330827

Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
nb=GaussianNB()
```

```
nb.fit(X_train,Y_train)
```

```
GaussianNB()
```

```
x_train_pred=nb.predict(X_train)
print(accuracy_score(Y_train,x_train_pred))
```

0.7849056603773585

```
x_test_pred=nb.predict(X_test)
bayes_data_accuracy=accuracy_score(Y_test,x_test_pred)
print(bayes_data_accuracy)
```

0.6992481203007519

```
print(classification_report(Y_test,x_test_pred))
```

	precision	recall	f1-score	support
0	0.77	0.80	0.78	90
1	0.54	0.49	0.51	43
accuracy			0.70	133
macro avg	0.65	0.64	0.65	133
weighted avg	0.69	0.70	0.70	133

```
print(confusion_matrix(Y_test,x_test_pred))
```

```
[[72 18]
 [22 21]]
```

KNN Classifier

```
from sklearn.neighbors import KNeighborsClassifier
kn=KNeighborsClassifier(n_neighbors=5,metric="euclidean",n_jobs=-1)
```

```
kn.fit(X_train,Y_train)
```

```
KNeighborsClassifier(metric='euclidean', n_jobs=-1)
```

```
xtrain_predict=kn.predict(X_train)
print(accuracy_score(Y_train,xtrain_predict))
```

0.8226415094339623

```
xtest_predict=kn.predict(X_test)
knn_data_accuracy=accuracy_score(Y_test,xtest_predict)
print(knn_data_accuracy)
```

0.7819548872180451

```
print(classification_report(Y_test,xtest_predict))
```

	precision	recall	f1-score	support
0	0.78	0.94	0.85	90
1	0.79	0.44	0.57	43
accuracy			0.78	133
macro avg	0.79	0.69	0.71	133
weighted avg	0.78	0.78	0.76	133

```
print(confusion_matrix(Y_test,xtest_predict))
```


```
[[85  5]
 [24 19]]
```

Comparing accuracy scores of 5 models

```
models=["SVM","Logistic","NaiveBayes","KNN"]
accuracies=[svm_data_accuracy,log_data_accuracy,bayes_data_accuracy,knn_data_accuracy]
barlist=plt.bar(models,accuracies,width=0.5,alpha=0.6)
```

```
barlist[1].set_color('r')
barlist[2].set_color('g')
barlist[3].set_color('y')
print(svm_data_accuracy,log_data_accuracy,bayes_data_accuracy,knn_data_accuracy)
```

```
0.7669172932330827 0.7593984962406015 0.6992481203007519
0.7819548872180451
```



Making a Predictive System

```
input_data=(4,110,92,0,37,6,0.191,30)
#changing input to numpy array
input_data_numpy_arr=np.asarray(input_data)

#reshape the array as we are predicting one instance
input_data_resaped=input_data_numpy_arr.reshape(1,-1)

#standardize input data
std_data=scaler.transform(input_data_resaped)
print(std_data)

prediction=classifier.predict(std_data)
print(prediction)

if prediction[0]:
    print("The person is diabetic")
else:
    print("The person is not diabetic")

[[ 0.04481216 -0.32339703  1.71581811 -1.41081572 -0.37717352 -4.04368046
 -0.94063914 -0.23068066]]
[0]
The person is not diabetic

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning:
X does not have valid feature names, but StandardScaler was fitted with
feature names
  warnings.warn(
```

```
import pickle

filename='diabetes_model.sav'
pickle.dump(classifier,open(filename,'wb'))
```

Heart Disease Prediction Model:

Importing Dependencies

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
```

Data Collection and Processing

```
#Loading csv data to pandas dataframe
df=pd.read_csv('/content/heart.csv')
heart_df=df.copy()
```

```
heart_df.head(5)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63.0	1	3	145	233	1	0	150.0	0	2.3	
1	37.0	1	2	130	250	0	1	187.0	0	3.5	
2	41.0	0	1	130	204	0	0	172.0	0	1.4	
3	56.0	1	1	120	236	0	1	178.0	0	0.8	
4	57.0	0	0	120	354	0	1	163.0	1	0.6	

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1

```
heart_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         276 non-null    float64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
```

```

5   fbs      303 non-null   int64
6   restecg  303 non-null   int64
7   thalach  280 non-null   float64
8   exang    303 non-null   int64
9   oldpeak  303 non-null   float64
10  slope    303 non-null   int64
11  ca       303 non-null   int64
12  thal     303 non-null   int64
13  target   303 non-null   int64
dtypes: float64(3), int64(11)
memory usage: 33.3 KB

```

```

#printing last 5 rows of dataset
heart_df.tail(5)

```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
298	57.0	0	0	140	241	0	1	NaN	1	0.2
299	45.0	1	3	110	264	0	1	132.0	0	1.2
300	68.0	1	0	144	193	1	1	141.0	0	3.4
301	57.0	1	0	130	131	0	1	115.0	1	1.2
302	57.0	0	1	130	236	0	0	174.0	0	0.0

	slope	ca	thal	target
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

```

#number of rows and columns in dataset
heart_df.shape

```

```

(303, 14)

```

```

heart_df.isnull().sum()

```


```

age      27
sex       0
cp        0
trestbps  0
chol      0
fbs       0
restecg   0
thalach   23
exang     0
oldpeak   0
slope     0
ca        0
thal      0
target    0
dtype: int64

```

```
#heatmap
import seaborn as sns
sns.heatmap(heart_df.isnull(),yticklabels=False,cbar=True,cmap="viridis")
```

<Axes: >



```
#distribution of data in age
fig,ax=plt.subplots(figsize=(8,8))
sns.distplot(heart_df.age)
```

<ipython-input-73-423799895cf4>:3: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(heart_df.age)
```

<Axes: xlabel='age', ylabel='Density'>



```
#distribution of data in thalach
fig,ax=plt.subplots(figsize=(8,8))
sns.distplot(heart_df.thalach)
```

<ipython-input-74-f84436d2239f>:3: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(heart_df.thalach)
```

<Axes: xlabel='thalach', ylabel='Density'>



#replace the missing values with median values


```
heart_df['age'].fillna(heart_df['age'].median(),inplace=True)
heart_df['thalach'].fillna(heart_df['thalach'].median(),inplace=True)
```

```
heart_df.isnull().sum()
```

```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
ca       0
thal     0
target   0
dtype: int64
```

```
sns.heatmap(heart_df.isnull(),yticklabels=False,cbar=True,cmap="viridis")
```


<Axes: >



```
f, axes = plt.subplots(figsize=(8, 8))
sns.distplot(heart_df.chol)
```

<ipython-input-78-051f31b8fda1>:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(heart_df.chol)
```

<Axes: xlabel='chol', ylabel='Density'>



```
def find_outliers(col):
    from scipy import stats
    z=np.abs(stats.zscore(col))
    idx_outliers=np.where(z>3,True,False)
    return pd.Series(idx_outliers,index=col.index)
idx=find_outliers(heart_df.chol)
```

idx

```
0      False
1      False
2      False
3      False
4      False
...
298    False
299    False
300    False
301    False
302    False
```

Length: 303, dtype: bool

```
display(heart_df.loc[idx==True].describe().round(3))
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach
count	4.000	4.0	4.000	4.000	4.000	4.00	4.0	4.000
mean	62.750	0.0	1.000	134.750	449.250	0.25	0.0	152.500
std	4.787	0.0	1.155	14.728	76.622	0.50	0.0	1.732
min	56.000	0.0	0.000	115.000	407.000	0.00	0.0	150.000
25%	61.250	0.0	0.000	129.250	408.500	0.00	0.0	152.250
50%	64.000	0.0	1.000	137.000	413.000	0.00	0.0	153.000
75%	65.500	0.0	2.000	142.500	453.750	0.25	0.0	153.250
max	67.000	0.0	2.000	150.000	564.000	1.00	0.0	154.000

	oldpeak	slope	ca	thal	target
count	4.000	4.00	4.000	4.00	4.000
mean	2.075	1.25	1.500	2.75	0.500
std	1.365	0.50	1.291	0.50	0.577
min	0.800	1.00	0.000	2.00	0.000
25%	1.400	1.00	0.750	2.75	0.000
50%	1.750	1.00	1.500	3.00	0.500
75%	2.425	1.25	2.250	3.00	1.000
max	4.000	2.00	3.000	3.00	1.000

```
display(heart_df.loc[idx==False].describe().round(3))
heart_df=heart_df.loc[idx==False]
```

	age	sex	cp	trestbps	chol	fbs	restecg	\
count	299.000	299.000	299.000	299.000	299.000	299.000	299.000	
mean	54.602	0.692	0.967	131.582	243.548	0.147	0.535	
std	8.743	0.462	1.032	17.590	45.858	0.355	0.526	
min	29.000	0.000	0.000	94.000	126.000	0.000	0.000	
25%	49.000	0.000	0.000	120.000	211.000	0.000	0.000	
50%	56.000	1.000	1.000	130.000	240.000	0.000	1.000	
75%	60.000	1.000	2.000	140.000	273.500	0.000	1.000	
max	77.000	1.000	3.000	200.000	394.000	1.000	2.000	

	thalach	exang	oldpeak	slope	ca	thal	target
count	299.000	299.000	299.000	299.000	299.000	299.000	299.000
mean	150.214	0.328	1.026	1.401	0.719	2.308	0.545
std	21.695	0.470	1.154	0.618	1.017	0.612	0.499
min	88.000	0.000	0.000	0.000	0.000	0.000	0.000
25%	138.500	0.000	0.000	1.000	0.000	2.000	0.000
50%	153.000	0.000	0.700	1.000	0.000	2.000	1.000
75%	165.000	1.000	1.600	2.000	1.000	3.000	1.000
max	202.000	1.000	6.200	2.000	4.000	3.000	1.000

```
f,axes=plt.subplots(figsize=(8,8))
sns.distplot(heart_df.chol)
```

<ipython-input-82-051f31b8fda1>:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(heart_df.chol)
```

<Axes: xlabel='chol', ylabel='Density'>



```
heart_df.corrwith(heart_df.target)
```

```
age      -0.224007
sex      -0.288805
cp        0.426488
trestbps -0.140696
chol     -0.104369
fbs      -0.037664
restecg   0.137884
thalach   0.405654
exang     -0.435335
oldpeak   -0.427715
slope     0.343767
ca        -0.384939
thal      -0.342352
target    1.000000
dtype: float64
```

```
heart_df.corr()
```

```

          age      sex      cp  trestbps      chol      fbs  \
age      1.000000 -0.120055 -0.053149  0.264521  0.191458  0.128176
sex     -0.120055  1.000000 -0.049753 -0.054249 -0.139533  0.051924
cp      -0.053149 -0.049753  1.000000  0.053367 -0.101672  0.086751
trestbps 0.264521 -0.054249  0.053367  1.000000  0.142659  0.176550
chol     0.191458 -0.139533 -0.101672  0.142659  1.000000  0.004715
fbs      0.128176  0.051924  0.086751  0.176550  0.004715  1.000000
restecg -0.082719 -0.079645  0.045442 -0.112884 -0.112560 -0.081747
thalach -0.338080 -0.024527  0.265942 -0.047277 -0.043502 -0.025900
exang    0.086527  0.141314 -0.392098  0.068965  0.092797  0.031749
oldpeak  0.195627  0.116748 -0.142571  0.188864  0.013709  0.012834
slope   -0.169671 -0.036135  0.115759 -0.123414  0.013755 -0.071281
ca       0.264302  0.136664 -0.171915  0.091869  0.052803  0.142793
thal     0.054926  0.228916 -0.158858  0.062779  0.066402 -0.023764
target  -0.224007 -0.288805  0.426488 -0.140696 -0.104369 -0.037664


          restecg  thalach      exang  oldpeak      slope      ca  \
age     -0.082719 -0.338080  0.086527  0.195627 -0.169671  0.264302
sex     -0.079645 -0.024527  0.141314  0.116748 -0.036135  0.136664
```

cp	0.045442	0.265942	-0.392098	-0.142571	0.115759	-0.171915
trestbps	-0.112884	-0.047277	0.068965	0.188864	-0.123414	0.091869
chol	-0.112560	-0.043502	0.092797	0.013709	0.013755	0.052803
fbs	-0.081747	-0.025900	0.031749	0.012834	-0.071281	0.142793
restecg	1.000000	0.084065	-0.073863	-0.047656	0.090725	-0.063052
thalach	0.084065	1.000000	-0.369432	-0.361457	0.352446	-0.216063
exang	-0.073863	-0.369432	1.000000	0.296586	-0.257863	0.115981
oldpeak	-0.047656	-0.361457	0.296586	1.000000	-0.577924	0.207042
slope	0.090725	0.352446	-0.257863	-0.577924	1.000000	-0.076253
ca	-0.063052	-0.216063	0.115981	0.207042	-0.076253	1.000000
thal	-0.002406	-0.109574	0.208045	0.199552	-0.096869	0.144648
target	0.137884	0.405654	-0.435335	-0.427715	0.343767	-0.384939

	thal	target
age	0.054926	-0.224007
sex	0.228916	-0.288805
cp	-0.158858	0.426488
trestbps	0.062779	-0.140696
chol	0.066402	-0.104369
fbs	-0.023764	-0.037664
restecg	-0.002406	0.137884
thalach	-0.109574	0.405654
exang	0.208045	-0.435335
oldpeak	0.199552	-0.427715
slope	-0.096869	0.343767
ca	0.144648	-0.384939
thal	1.000000	-0.342352
target	-0.342352	1.000000

```
plt.figure(figsize = (10, 5))
sns.heatmap(heart_df.corr(), annot = True, annot_kws = {"size":10})
```

<Axes: >



#statistical measure about the data

heart_df.describe()

	age	sex	cp	trestbps	chol
count	299.000000	299.000000	299.000000	299.000000	299.000000
mean	54.602007	0.692308	0.966555	131.581940	243.548495
std	8.743481	0.462312	1.032469	17.589726	45.857602
min	29.000000	0.000000	0.000000	94.000000	126.000000
25%	49.000000	0.000000	0.000000	120.000000	211.000000
50%	56.000000	1.000000	1.000000	130.000000	240.000000
75%	60.000000	1.000000	2.000000	140.000000	273.500000
max	77.000000	1.000000	3.000000	200.000000	394.000000

	restecg	thalach	exang	oldpeak	slope
count	299.000000	299.000000	299.000000	299.000000	299.000000
mean	0.535117	150.214047	0.327759	1.025753	1.401338
std	0.525782	21.694671	0.470183	1.154495	0.618071
min	0.000000	88.000000	0.000000	0.000000	0.000000
25%	0.000000	138.500000	0.000000	0.000000	1.000000
50%	1.000000	153.000000	0.000000	0.700000	1.000000
75%	1.000000	165.000000	1.000000	1.600000	2.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000

	thal	target
count	299.000000	299.000000
mean	2.307692	0.545151
std	0.612214	0.498792
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

#checking the distribution of target variable

heart_df.target.value_counts()

```
1    163
0    136
Name: target, dtype: int64
```

1--->Defective Heart 0--->Healthy Heart

Splitting the features and target

```
X=heart_df.drop('target',axis=1)
Y=heart_df['target']
```

```
print(X)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak
\										
0	63.0	1	3	145	233	1	0	150.0	0	2.3
1	37.0	1	2	130	250	0	1	187.0	0	3.5
2	41.0	0	1	130	204	0	0	172.0	0	1.4
3	56.0	1	1	120	236	0	1	178.0	0	0.8
4	57.0	0	0	120	354	0	1	163.0	1	0.6
..
298	57.0	0	0	140	241	0	1	153.0	1	0.2
299	45.0	1	3	110	264	0	1	132.0	0	1.2
300	68.0	1	0	144	193	1	1	141.0	0	3.4
301	57.0	1	0	130	131	0	1	115.0	1	1.2
302	57.0	0	1	130	236	0	0	174.0	0	0.0

	slope	ca	thal
0	0	0	1
1	0	0	2
2	2	0	2
3	2	0	2
4	2	0	2
..
298	1	0	3
299	1	0	3
300	1	2	3
301	1	1	3
302	1	1	2

```
[299 rows x 13 columns]
```

```
print(Y)
```

```
0    1
1    1
2    1
3    1
4    1
..
298  0
299  0
300  0
301  0
```


302 0

Name: target, Length: 299, dtype: int64

Splitting the data into training and test data

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=2,stratify=Y)
```

```
print(X.shape,X_train.shape,X_test.shape)
```

```
(299, 13) (239, 13) (60, 13)
```

Model Training

```
#Logistic Regression
```

```
clf=LogisticRegression()
```

```
#training model with training data
```

```
clf.fit(X_train,Y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
    n_iter_i = _check_optimize_result(
```

```
LogisticRegression())
```

Model evaluation

```
#accuracy on training data
```

```
X_train_prediction=clf.predict(X_train)
```

```
training_data_accuracy=accuracy_score(Y_train,X_train_prediction)
```

```
print("Accuracy on Training data ",training_data_accuracy)
```

```
Accuracy on Training data   0.8410041841004184
```

```
#accuracy on test data
```

```
X_test_prediction=clf.predict(X_test)
```

```
log_data_accuracy=accuracy_score(Y_test,X_test_prediction)
```

```
print("Accuracy on Test data ",log_data_accuracy)
```

```
print(classification_report(Y_test,X_test_prediction))
```

```
print(confusion_matrix(Y_test,X_test_prediction))
```

```
Accuracy on Test data   0.85
```

```
          precision        recall    f1-score     support
```

```
          0            0.82        0.85        0.84          27
```

	1	0.88	0.85	0.86	33
accuracy				0.85	60
macro avg		0.85	0.85	0.85	60
weighted avg		0.85	0.85	0.85	60

```
[[23  4]
 [ 5 28]]
```

SVM

```
from sklearn import svm
classifier=svm.SVC(kernel="linear")
classifier.fit(X_train,Y_train)

SVC(kernel='linear')

#accuracy score on the training data
svc_x_train_prediction=classifier.predict(X_train)
svc_training_data_accuracy=accuracy_score(svc_x_train_prediction,Y_train)

print(svc_training_data_accuracy)

0.8451882845188284

svc_x_test_prediction=classifier.predict(X_test)
svm_data_accuracy=accuracy_score(svc_x_test_prediction,Y_test)
print(svm_data_accuracy)

0.8166666666666667

print(classification_report(Y_test,svc_x_test_prediction))
print(confusion_matrix(Y_test,svc_x_test_prediction))
```

		precision	recall	f1-score	support
	0	0.86	0.70	0.78	27
	1	0.79	0.91	0.85	33
accuracy				0.82	60
macro avg		0.83	0.81	0.81	60
weighted avg		0.82	0.82	0.81	60

```
[[19  8]
 [ 3 30]]
```

Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
rclf=RandomForestClassifier()

rclf.fit(X_train,Y_train)

RandomForestClassifier()

train_predictions=rclf.predict(X_train)
print(accuracy_score(Y_train,train_predictions))
```

1.0

```
test_predictions=rclf.predict(X_test)
random_data_accuracy=accuracy_score(Y_test,test_predictions)
print(random_data_accuracy)
```

0.85

```
print(classification_report(Y_test,test_predictions))
print(confusion_matrix(Y_test,test_predictions))
```

	precision	recall	f1-score	support
0	0.88	0.78	0.82	27
1	0.83	0.91	0.87	33
accuracy			0.85	60
macro avg	0.85	0.84	0.85	60
weighted avg	0.85	0.85	0.85	60

```
[[21  6]
 [ 3 30]]
```

Naive Bayes

```
from sklearn.naive_bayes import GaussianNB
nb=GaussianNB()
```

```
nb.fit(X_train,Y_train)
```

```
GaussianNB()
```

```
x_train_pred=nb.predict(X_train)
print(accuracy_score(Y_train,x_train_pred))
```

0.8326359832635983

```
x_test_pred=nb.predict(X_test)
bayes_data_accuracy=accuracy_score(Y_test,x_test_pred)
print(bayes_data_accuracy)
```

0.8166666666666667

```
print(classification_report(Y_test,x_test_pred))
print(confusion_matrix(Y_test,x_test_pred))
```

	precision	recall	f1-score	support
0	0.81	0.78	0.79	27
1	0.82	0.85	0.84	33
accuracy			0.82	60
macro avg	0.82	0.81	0.81	60
weighted avg	0.82	0.82	0.82	60

```
[[21 6]
 [ 5 28]]
```

KNN

```
from sklearn.neighbors import KNeighborsClassifier
kn=KNeighborsClassifier(n_neighbors=5,metric="euclidean",n_jobs=-1)
```

```
kn.fit(X_train,Y_train)
```

```
KNeighborsClassifier(metric='euclidean', n_jobs=-1)
```

```
xtrain_predict=kn.predict(X_train)
print(accuracy_score(Y_train,xtrain_predict))
```

```
0.7531380753138075
```

```
xtest_predict=kn.predict(X_test)
knn_data_accuracy=accuracy_score(Y_test,xtest_predict)
print(knn_data_accuracy)
```

```
0.7166666666666667
```

```
print(classification_report(Y_test,xtest_predict))
```

	precision	recall	f1-score	support
0	0.73	0.59	0.65	27
1	0.71	0.82	0.76	33
accuracy			0.72	60
macro avg	0.72	0.71	0.71	60
weighted avg	0.72	0.72	0.71	60

```
print(confusion_matrix(Y_test,xtest_predict))
```


```
[[16 11]
 [ 6 27]]
```

Comparing accuracies of 5 models

```
models=["SVM","Logistic","NaiveBayes","KNN"]
accuracies=[svm_data_accuracy,log_data_accuracy,bayes_data_accuracy,knn_data_accuracy]
barlist=plt.bar(models,accuracies,width=0.5,alpha=0.6)
```

```
barlist[1].set_color('r')
barlist[2].set_color('g')
barlist[3].set_color('y')
print(svm_data_accuracy,log_data_accuracy,bayes_data_accuracy,knn_data_accuracy)
```

```
0.8166666666666667 0.85 0.8166666666666667 0.7166666666666667
```



Building a predictive system

```
input_data=(62,0,0,140,268,0,0,160,0,3.6,0,2,2)
#change input data to numpy array
input_data_as_numpy_array=np.asarray(input_data)
#reshape the numpy array as we are predicting for only one instance
input_data_reshaped=input_data_as_numpy_array.reshape(1,-1)
prediction=clf.predict(input_data_reshaped)
print(prediction)
if prediction:
    print("Person has heart disease")
else:
    print("Person doesn't have heart disease")
```

[0]

Person doesn't have heart disease

```
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning:
X does not have valid feature names, but LogisticRegression was fitted
with feature names
  warnings.warn(
```

```
import pickle
filename='heart_model.sav'
pickle.dump(clf,open(filename,'wb'))
```

Making Multiple Disease Prediction System using streamlit:

```
import pickle
import streamlit as st
from streamlit_option_menu import option_menu

#loading saved models
diabetes_model=pickle.load(open("C:/Users/taman/Desktop/Multiple Disease Prediction/diabetes_model.sav","rb"))
heart_model=pickle.load(open("C:/Users/taman/Desktop/Multiple Disease Prediction/heart_model.sav","rb"))
parkinson_model=pickle.load(open("C:/Users/taman/Desktop/Multiple Disease Prediction/parkinson_model.sav","rb"))

#side bar/navigation
with st.sidebar:
    selected=option_menu('Multiple Disease Prediction System',['Diabetes Prediction','Heart Disease Prediction'],icons=['activity','heart'],)

#Diabetes Prediction System
if (selected=='Diabetes Prediction'):
    #page title
    st.title('Diabetes Prediction using ML')

    Pregnancies=st.text_input('Number of Pregnancies')
    Glucose=st.text_input('Glucose Level')
    BloodPressure=st.text_input('Blood Pressure Value')
    SkinThickness=st.text_input('Skin Thickness Value')
    Insulin=st.text_input('Insulin Level')
    BMI=st.text_input('BMI value')
    DiabetesPedigreeFunction=st.text_input('Diabetes Pedigree Function Value')
    Age=st.text_input('Age of the person')

    #code for Prediction
    diab_diagnosis=''

    #creating a button
    if st.button('Diabetes Test Result'):
        diab_prediction=diabetes_model.predict([[Pregnancies,Glucose,BloodPressure,SkinThickness,Insulin,BMI,DiabetesPedigreeFunction,Age]])
        if diab_prediction:
            diab_diagnosis='The person is Diabetic'
        else:
            diab_diagnosis='The person is not diabetic'

    diab_diagnosis='The person is not diabetic'
    st.success(diab_diagnosis)

#Heart Disease Prediction
if (selected=='Heart Disease Prediction'):
    #page title
    st.title('Heart Disease Prediction using ML')

    # age=st.text_input('Age')
    # sex=st.text_input('Sex (0 = Female; 1 = Male)')
    # chestpain=st.text_input('Chest Pain 0-3')
    #restingbp=st.text_input('Resting Blood Pressure')
    #cholesterol=st.text_input('Serum cholesterol in mg/dl')
    #bloodsugar=st.text_input('fasting blood sugar>120 mg/dl (1 = True; 0 = False)')
    #ecardio=st.text_input('Resting electrocardiographic results')
    #maxheartrate=st.text_input('Maximum heart rate achieved')
    #angina=st.text_input('Exercise induced angina (1 = yes; 0 = no)')
    #oldpeak=st.text_input('ST depression induced by exercise relative to rest')
    #slope=st.text_input('The slope of the peak exercise ST segment')
    #number=st.text_input('Number of major vessels(0-3) colored by flourosopy')
    #thal=st.text_input('thal(0 = normal; 1 = fixed defect; 2 = reversable effect)')

    #heart_diag=''

    if st.button('Heart Disease Test Result'):
        # heart_prediction=heart_model.predict([[age,sex,chestpain,restingbp,cholesterol,bloodsugar,ecardio,maxheartrate,angina,oldpeak,slope]])
        # if heart_prediction:
        #     heart_diag='The person has a Defective Heart'
        # else:
        #     heart_diag='The person has a Healthy Heart'
        st.success(heart_diag)

    # Heart Disease Prediction Page
if (selected == 'Heart Disease Prediction'):
    # page title
    st.title('Heart Disease Prediction using ML')

    col1, col2, col3 = st.columns(3)
```

```

with col1:
    age = st.number_input('Age')

with col2:
    sex = st.number_input('Sex')

with col3:
    cp = st.number_input('Chest Pain types')

with col1:
    trestbps = st.number_input('Resting Blood Pressure')

with col2:
    chol = st.number_input('Serum Cholestoral in mg/dl')

with col3:
    fbs = st.number_input('Fasting Blood Sugar > 120 mg/dl')

with col1:
    restecg = st.number_input('Resting Electrocardiographic results')
with col2:
    thalach = st.number_input('Maximum Heart Rate achieved')

with col3:
    exang = st.number_input('Exercise Induced Angina')

with col1:
    oldpeak = st.number_input('ST depression induced by exercise')

with col2:
    slope = st.number_input('Slope of the peak exercise ST segment')

with col3:
    ca = st.number_input('Major vessels colored by flourosopy')

with col1:
    thal = st.number_input('thal: 0 = normal; 1 = fixed defect; 2 = reversable defect')

```

```

# code for Prediction
heart_diagnosis = ''

# creating a button for Prediction
if st.button('Heart Disease Test Result'):
    heart_prediction = heart_model.predict([[age, sex, cp, trestbps, chol, fbs, restecg, thalach, exang, oldpeak, slope, ca, thal]])

    if (heart_prediction[0] == 1):
        heart_diagnosis = 'The person is having heart disease'
    else:
        heart_diagnosis = 'The person does not have any heart disease'

st.success(heart_diagnosis)

```

FUTURE SCOPE OF IMPROVEMENT

ML models heavily rely on high-quality and diverse data for accurate predictions. Future improvements can focus on collecting larger and more comprehensive datasets that include a wider range of patient demographics, genetic information, lifestyle factors, and environmental data. Incorporating real-time data from wearable devices and electronic health records can also provide valuable insights.

Deep learning algorithms, such as convolutional neural networks (CNNs) and recurrent neural networks (RNNs), have shown great potential in healthcare applications.

ML can be leveraged to develop personalized disease prediction models that consider an individual's unique characteristics, such as genetic makeup, medical history, lifestyle, and environmental factors. By tailoring predictions and treatment recommendations to specific patients, healthcare providers can offer more targeted and effective interventions.

CERTIFICATE

This is to certify that Miss Tamanna Sharma of Lovely Professional University, registration number: 12110221, has successfully completed a project on *Multiple Disease prediction* using *Machine Learning with Python* under the guidance of Prof. Arnab Chakraborty.

Prof. Arnab Chakraborty

Globsyn Finishing School

